




THE MANUAL for the

SYSTEMATIC INVESTOR

*A guide to put capital to work in
systematic, rules-based investment*

JOHN PANAYOTARAS

THE MANUAL FOR THE SYSTEMATIC INVESTOR

*A guide for the serious investor who wants to put capital
to work in a systematic, rules-based investment strategy.*

FOR HIGH-NET-WORTH INVESTORS

John Panayotaras

INSIGHT LLC

Copyright © 2026 Insight LLC

All rights reserved.

Written by John Panayotaras.

Published by Insight LLC.

No part of this publication may be reproduced, distributed, or transmitted in any form or by any means, including photocopying, recording, or other electronic or mechanical methods, without the prior written permission of the publisher, except in the case of brief quotations embodied in critical reviews and certain other non-commercial uses permitted by copyright law.

For permission requests, contact: Insight LLC.

This book is intended for informational and educational purposes only. It does not constitute investment advice, financial advice, legal advice, or any other form of professional advice. See the full disclaimer inside.

First Edition, 2026

Printed and distributed by Insight LLC

For those who ask the harder questions.

Contents

Disclaimer — A Few Words Before We Start	vii
Preface — A Note Before We Begin	ix
1 The Language of the Game	1
2 The Mutual Fund Story: A Cautionary Tale	7
3 The Origins of Systematic Investing: Where the Smart Money Went	13
4 Building a Systematic Strategy — Part One: The Foundation	19
5 Building a Systematic Strategy — Part Two: Making It Real	23
6 Artificial Intelligence in Systematic Strategies: The Truth	29
7 Risk Management: The Only Thing That Keeps You Alive	33
8 The S&P 500 Fallacy: Why America Is Not the World	39
9 Your Portfolio: Choosing and Sizing a Systematic Strategy	45
10 Epilogue: Why Systematic Matters — and Why the Story Is Bigger Than That	53

Key Ideas, Illustrated	57
Key Terms, Simply Explained	61
About the Author	69

A Few Words Before We Start

Let me be upfront with you about something.

This book is not investment advice. It is not a recommendation to buy or sell anything, and nothing in these pages should be taken as a substitute for speaking with a qualified financial adviser, a tax professional, or a lawyer who knows your specific situation. Every investor is different. Every portfolio is different. Every set of circumstances is different. What works for one person may be entirely wrong for another.

The author does not manage other people's money, has no financial interest in any fund or product mentioned, and receives nothing from anyone whose strategies are discussed here. Any reference to academic research or specific investment approaches is purely for educational context — not an endorsement.

Here is what this book actually is: a long, honest conversation about a corner of the investment world that most people find confusing, and that the financial industry would generally prefer you remain confused about. It is the kind of conversation you might have with a knowledgeable friend at dinner — someone willing to tell you the truth but also perfectly clear that they are not your broker, your accountant, or your legal counsel.

Past performance, as you have almost certainly heard, does not guarantee future results. All investments carry risk. Some carry more than you might expect. A few carry risks that even sophisticated investors do not fully appreciate until it is too late.

Read this book with curiosity and appropriate scepticism. Then go talk to the right professionals.

A Note Before We Begin

If you have accumulated significant wealth — whether through building a business, rising through a profession, or inheriting a family fortune — you have almost certainly encountered a parade of people who want to help you manage it. Bankers, advisors, brokers, consultants, wealth managers, and “alternative investment specialists” have all knocked on your door with a smile and a pitch deck.

This book is not a pitch. It does not end with a product to sell or a fund to subscribe to. Instead, it is a guide to understanding a specific corner of the investment world that is often described with impressive-sounding language but rarely explained clearly: systematic investing.

Systematic investing — sometimes called quantitative investing, algorithmic trading, or rules-based investing — is a discipline that has, for decades, quietly generated some of the most extraordinary returns in financial history. And yet, it remains poorly understood by most private investors, even very sophisticated ones.

Why? Partly because the practitioners prefer it that way. Partly because the mathematics involved can be genuinely complex. And partly because, frankly, the financial industry has no strong incentive to explain it to you clearly — they would rather keep charging fees for services that could be replaced by something better.

This book will change that. Chapter by chapter, we will walk through what systematic investing is, where it came from, how it works, what its genuine strengths are, and — perhaps most impor-

tantly — how you, as a high-net-worth investor, might intelligently incorporate it into your portfolio.

We will keep things plain. When a technical term is unavoidable, we will explain it in plain English. When the industry likes to dress up a simple idea in complexity, we will undress it. And occasionally, when the situation calls for it, we will allow ourselves a dry smile at how predictably absurd this business can be.

Let us begin.

CHAPTER ONE

The Language of the Game


Before we can talk about investing intelligently, we need to agree on what certain words actually mean. The investment industry has a remarkable talent for taking simple concepts and wrapping them in jargon until even the practitioners are not entirely sure what they are talking about. So let us start by defining our terms with precision and, wherever possible, with honesty.

What We Mean by “Systematic”

When we say that an investment strategy is systematic, we mean that it follows a fixed set of rules. A systematic strategy does not rely on a person’s gut feeling, their reading of the morning newspaper, their network of contacts, or their ability to charm a CEO over lunch. Instead, it uses a predefined process — a set of instructions — to make investment decisions.

Think of it like a recipe. A good chef does not randomly add ingredients and hope for the best. They follow a tested recipe, adjusting quantities in a disciplined way, and producing a result that can be replicated. A systematic strategy is the investment equivalent of a kitchen that has replaced guesswork with a well-tested process.

Plate I



The Discipline of Rules

—◇—
a strategy measured the same way, every time

This might sound simple, but it is actually quite profound. Most of us make financial decisions the way we make most life decisions: based on how we feel, what we recently heard, and what stories our minds are currently telling us. A systematic strategy is designed specifically to eliminate that kind of decision-making.

Systematic strategies can be built around many different types of ideas: buying assets that have risen in price (momentum), buying assets that are cheap relative to their earnings (value), or trading based on changes in interest rates, currencies, commodities, or hundreds of other factors. What makes them all systematic is not what they invest in, but how: according to clear, testable, repeatable rules.

It is worth noting explicitly: systematic does not mean high-frequency or short-term. A systematic strategy might hold positions

for weeks, months, or even years. Some of the most successful systematic funds in the world take a view on interest rates or currency trends and hold their positions patiently over many months. Speed is not the defining characteristic — discipline is.

What We Mean by “Algorithmic Trading”

Algorithmic trading is the execution of trades using computer programs rather than human traders. In its most basic form, an algorithm is simply a set of instructions a computer follows: if the price of Asset X rises above a certain level, buy a certain quantity of Asset X. In its most complex form, a trading algorithm might process millions of data points per second and adjust hundreds of positions simultaneously.

Much of what appears in the media when algorithmic trading is mentioned refers to high-frequency trading (HFT) — strategies that exploit tiny price differences over milliseconds, sometimes microseconds. These strategies bear almost no resemblance to what most serious systematic investors do.

When we talk about systematic investing in this book, we are not talking about HFT. We are not talking about computers racing each other across fiber-optic cables at the speed of light to snatch fractions of a cent per share. We are talking about disciplined, rule-based investment strategies that operate on timescales of days, weeks, months, or longer. The computer is a tool for consistency and scale, not a device for playing speed games.

What We Mean by “Discretionary”

Discretionary investing is the opposite of systematic investing. In a discretionary approach, a human being — an analyst, a portfolio manager, a fund manager — makes investment decisions based on their judgment. They read company reports, talk to management teams, analyse economic data, and ultimately decide, using their own intelligence and experience, what to buy and sell.

This is not a bad thing. Some of the greatest investors in history have been discretionary investors of extraordinary skill. Discretionary investing, done well, is a genuine art. The problem — as we will explore in Chapter Two — is that discretionary investing done poorly looks almost identical to discretionary investing done well, at least in the short term. It is very difficult to distinguish genuine skill from luck in financial markets, and the financial industry has not always been honest about this.

The discretionary world has its heroes. It also has many well-dressed, highly confident practitioners who are, in the kindest possible terms, not quite as good as they believe themselves to be.

Passive Investing vs. Active Management

The simplest distinction in investing is between passive and active. Let us be very clear about what each means.

Passive investing means buying a basket of investments that tracks a market — usually an index like the S&P 500 — and holding it. You are not trying to beat the market. You are trying to be the market, minus a small cost. A passive investor buys an ETF (Exchange-Traded Fund) that tracks the S&P 500 and goes home. The market goes up, they go up. The market goes down, they go down. No analysis. No effort. No drama.

Active management means trying to do better than the market. An active manager — whether discretionary or systematic — is attempting to generate returns that exceed what a simple passive index would deliver. In industry terminology, this excess return is called alpha. (Beta is simply your exposure to market returns — the part that goes up and down with the overall market.)

Passive investing is cheap, simple, and — here is the uncomfortable truth that most of the financial industry would prefer you not to think too hard about — often more successful than expensive active management. We will spend considerable time on this.

But the world is more nuanced than passive good, active bad. There are forms of active management, including certain system-

atic approaches, that have generated genuine, sustained alpha over many years. Finding them, evaluating them, and allocating to them intelligently is what much of this book is about.

Now that we have our vocabulary in order, let us go back in time to understand how we got here.

CHAPTER TWO

The Mutual Fund Story: A Cautionary Tale

To understand where systematic investing came from, it helps to understand what it was reacting against. And the clearest starting point is the mutual fund industry — one of the most successful marketing operations in financial history.

What Was the Mutual Fund?

A mutual fund is, at its core, a simple idea: pool money from many investors, hire a professional manager to invest it, and share the profits (and losses) proportionally. The concept became popular in the United States in the 1950s and 1960s, and it solved a genuine problem. Individual investors did not have the time, knowledge, or capital to build diversified portfolios of individual stocks. A mutual fund gave them access to professional management and broad diversification in exchange for a management fee.

This was a reasonable trade when it worked. And for several decades, it appeared to work spectacularly well.

Why It Was Easy to Make Money for a Long Time

Here is something the mutual fund industry's marketing materials tend not to highlight: a significant portion of the strong returns that mutual funds delivered from the 1950s through the 1990s had less to do with managerial brilliance and more to do with a remarkable macroeconomic backdrop.

The post-World War II period was one of the most extraordinary periods of economic growth in human history. Populations were growing. Technology was advancing. Companies that had been devastated by the war were rebuilding. Inflation, after some volatility, eventually fell dramatically through the 1980s, which caused bond prices and stock valuations to rise substantially. A manager who simply rode these waves looked like a genius.

Then there was the expansion of the middle class, the rise of retirement savings vehicles, and — critically — the sheer volume of retail investors entering the stock market for the first time, often without sophisticated analysis. In such an environment, a competent, research-driven fund manager could genuinely add value by identifying good companies before the less-informed crowd noticed them.

The game was relatively easier, and the field was less crowded. Smart people with access to company research could, and did, outperform. The performance records looked excellent, the fees seemed justified, and mutual funds became embedded in the financial plans of millions of households.

What the Evidence Actually Shows

The trouble began when researchers started asking a simple question: do professional fund managers actually outperform a simple passive index, after adjusting for their fees? The answer, accumulated across decades of careful academic study, was deeply uncomfortable.


The majority of actively managed mutual funds underperform their benchmark index after fees. Not marginally. Not occasionally.

Systematically, persistently, and across most time periods and most categories studied. Study after study, across markets in the United States, Europe, Asia, and elsewhere, arrived at the same basic conclusion: on average, active fund managers subtract value rather than add it, once costs are properly accounted for.

How is this possible? These are smart, hardworking, highly educated people. They have access to enormous amounts of information. They spend their professional lives analysing companies and markets. How can they consistently underperform a dumb index that just holds everything proportionally?

The answer involves several forces. First, markets are competitive. When smart, well-resourced people compete intensely with each other, the advantages tend to cancel out. Second, and critically, fees are real and compounding. A fund charging 1.5% per year in fees needs to generate at least 1.5% per year of additional return before even breaking even against a free index. Over twenty years, 1.5% per year compounded means a fund must outperform by more than 30% in total just to keep you at the same place. Third, there is a mathematical reality: the average active manager must, by definition, produce average returns before fees. After fees, the average active manager produces below-average returns.

Plate II



The Quiet Tax

a small fee, compounded into a fortune lost

How Quantitative Finance Was Born from Disappointment

The evidence against active discretionary management was accumulating through the 1970s and 1980s, and a generation of mathematicians, physicists, and statisticians began to look at financial markets with fresh eyes. Rather than asking which stocks should I pick, they began asking: Are there patterns in price data that can be exploited systematically? Are there risk factors that consistently explain returns? Can a disciplined, rule-based approach capture returns that human managers consistently fail to capture?

From this intellectual ferment emerged quantitative finance, and from quantitative finance emerged the first generation of systematic

investment funds. These were not built on tips, relationships, or intuition. They were built on data, statistics, and rigorous testing. The results, as we will see in the next chapter, were remarkable.

CHAPTER THREE

The Origins of Systematic Investing: Where the Smart Money Went

The story of systematic investing is a story about scientists entering finance and refusing to accept conventional wisdom without evidence. It begins in the 1970s and accelerates dramatically through the 1980s and 1990s. Let us trace the arc.

How It Started — and Why

The intellectual seeds were planted by a generation of academics who approached financial markets as a scientific problem. Eugene Fama's work on market efficiency established the intellectual framework: if all available information is already priced into markets, then no amount of analysis should allow you to consistently outperform. This was deeply provocative to an industry built on the belief that smart analysts could always find the angle.

At the same time, practitioners like Ed Thorp — a mathematics professor who had already beaten casinos at blackjack using systematic card-counting — were applying quantitative methods to financial markets. Thorp's insight was elegant: financial markets,

like casinos, have statistical regularities that can be exploited if you are disciplined enough to follow the rules without emotion.

Through the 1970s and 1980s, a small number of firms began building systematic trading operations. Commodities Corporation, based in Princeton, began funding quantitative traders who used statistical models to trade futures markets. From this stable emerged many of the legendary names in systematic trading, including Paul Tudor Jones and Bruce Kovner — though both evolved toward more discretionary approaches. The purely quantitative tradition continued in firms like D.E. Shaw and Renaissance Technologies.

Renaissance Technologies, founded by mathematician James Simons, became perhaps the most famous example of systematic investing done at the highest level. Its Medallion Fund — available only to employees — generated average annual returns of roughly 66% before fees over three decades. Even after the fund's famously steep fees, returns were extraordinary. The fund employed mathematicians, physicists, cryptographers, and computer scientists. It did not employ traditional financial analysts. Its models were, and remain, highly secretive.

The Performance Data of Systematic Funds

The broader systematic investing industry — which includes managed futures funds (also known as Commodity Trading Advisors, or CTAs), quantitative equity funds, and systematic macro funds — has compiled a performance record that, in aggregate, looks very different from the mutual fund world.

The managed futures sector in particular has demonstrated two characteristics that make it extraordinarily valuable to institutional investors. First, it has generated positive returns over the long run. Second, and more importantly, it has tended to perform well during periods when stock markets perform badly. This is the investment equivalent of an umbrella that actually works.

During the global financial crisis of 2008, when the S&P 500 fell approximately 37%, many managed futures funds generated

strongly positive returns. During the dot-com bust of 2000 to 2002, while traditional equity funds were suffering catastrophic losses, managed futures funds were often making money. This is not a coincidence — it is the result of strategies that can systematically profit from both rising and falling markets.

The Problem of Access

Here is the catch, and it is a significant one: the best systematic investment funds are, in many cases, either closed to new investors entirely or accessible only to very large institutions under very demanding conditions.

Renaissance Technologies' Medallion Fund, as mentioned, is closed to outside investors. AHL, the systematic trading arm of Man Group — one of the world's largest systematic managers — has minimum investment thresholds appropriate only for large institutions, and its fee structure reflects the premium nature of its offering. Two Sigma, Winton, and similar firms operate with sophisticated institutional clients as their audience.

The fees in this world can be substantial. The traditional hedge fund fee structure of two and twenty — meaning a 2% annual management fee plus 20% of profits — is common in the systematic hedge fund space. Some of the best firms charge more. These fees are not outrageous if the underlying performance justifies them, but they require careful analysis before committing capital.

For high-net-worth investors, access has traditionally been difficult but is becoming more available. Managed futures mutual funds and ETFs now exist in the retail space, though they do not always capture the full return profile of the institutional versions. The private wealth space has seen increasing availability of systematic strategies through managed accounts and dedicated alternative investment vehicles.

Is Systematic Investing Bad Because It Has No Economic Analysis?

This is one of the most common objections raised by traditional investors, and it deserves a direct answer: no. The absence of traditional economic analysis does not make systematic investing economically meaningless.

Here is why. Financial markets serve essential economic functions: they aggregate information about the future value of assets, they provide liquidity to buyers and sellers, and they transfer risk from those who do not want it to those who do. A systematic strategy that buys assets with upward price momentum is not operating in an economic vacuum — it is participating in the process of incorporating new information into prices. A systematic strategy that takes the other side of hedgers in commodity futures markets is providing a genuine economic service: liquidity and risk transfer.

Furthermore, systematic strategies can invest in exactly the same assets as any other investor: stocks, bonds, currencies, commodities, and so on. A systematic strategy that identifies high-quality, low-debt, growing businesses and holds them patiently is doing something that even the most traditional value investor would recognise as economically sensible.

Is Discretionary Management a Fraud?

No. It is not a fraud. But it is rarer as a source of genuine skill than the industry would have you believe.

There are discretionary investors of extraordinary ability — people who can consistently identify mispriced assets before the market corrects the pricing. Warren Buffett's career is the most famous example. His multi-decade record of outperformance is statistically almost impossible to attribute to luck. He is genuinely exceptional.

The problem is that there are many people in the discretionary investment world who believe, sincerely, that they are in the same category. Some of them are wrong. The honest assessment is this: genuine discretionary skill exists, but it is rare, difficult to identify in

advance, and increasingly crowded out by the volume and speed of information available to all market participants.

The relevant question for a serious investor is not whether discretionary management can work in theory. It is whether you can reliably identify the handful of managers who are truly skilled, distinguish them from the many who are not, and access their services at a price that leaves a meaningful return for you. That is a demanding challenge, and one we will return to.

CHAPTER FOUR

Building a Systematic Strategy — Part One: The Foundation

Now we begin to peel back the curtain. How is a systematic investment strategy actually built? What does the process look like from the beginning? This chapter and the next will walk you through it in plain language, without the mathematics that typically surrounds these explanations. Think of it as watching a restaurant kitchen from a window — you will see how the food is made without needing to be a chef yourself.

The Initial Idea: Where Does a Strategy Begin?

Every systematic strategy starts with an idea. Not an arbitrary idea — a specific, testable hypothesis about how financial markets behave. This hypothesis is usually rooted in one of the following sources:

Market phenomena: Observable patterns in how prices move. For example: assets that have risen in price over the past year tend, on average, to continue rising over the next few months. This is the momentum phenomenon, and it has been documented across stocks,

bonds, commodities, and currencies across many decades and many markets.

Risk factor theory: The idea that certain types of assets carry higher risk, and that bearing that risk deserves compensation in the form of higher returns. Value stocks — companies whose share prices are low relative to what the company actually earns or owns — have historically outperformed over long periods. The theory is that these companies are genuinely riskier or less popular, and investors demand higher expected returns to hold them.

Structural market features: Some participants in financial markets are not trying to maximise returns. A wheat farmer wants to sell futures contracts to lock in a price for their harvest — they are protecting themselves against a bad price, not speculating. A systematic strategy that provides this service can earn a structural return over time.

The hypothesis must be specific and, critically, it must be rooted in a reasonable economic logic. A strategy that has worked historically but whose returns cannot be explained by any plausible mechanism is a red flag — it may be the result of data mining rather than genuine insight. We will return to this distinction.

Historical Data: The Raw Material — and the Danger

Once an idea exists, the strategy developer turns to historical data to examine whether the idea has merit. Historical price data, corporate financial data, economic data, and many other information sources are used to build and test systematic strategies.

Historical data is essential. It is also dangerous if misused.

The most important risk in this phase is survivorship bias. If you test a stock-picking strategy using only the companies that exist in today's stock market, you will get misleadingly optimistic results — because you are excluding all the companies that went bankrupt, merged, or otherwise disappeared from the data. A strategy that looks excellent on a survivorship-biased dataset may look very different when tested correctly.

A second risk is look-ahead bias: accidentally using information that would not have been available at the time the strategy would have been trading. For example, if your strategy uses quarterly earnings data, it must only use that data after it has been publicly released — not on the date of the quarter ended.

A third risk, which we will discuss in depth, is the most insidious: overfitting. This is the danger of finding patterns that are accidents of the specific historical data rather than genuine, persistent market phenomena. This is the systematic investor's equivalent of mistaking a lucky streak for skill.

Strategy Development and Programming

With a hypothesis and clean historical data, the developer begins constructing the strategy. This involves defining the precise rules: how to identify the signal (the condition that triggers a trade), how much to buy or sell, and under what conditions to exit a position.

This process is iterative. A developer might start with a simple version of their idea — perhaps just buying the top 20% of stocks ranked by twelve-month price momentum — and then examine the results carefully. Do the returns make economic sense? Are they concentrated in a particular period, sector, or market condition? Are transaction costs — the cost of actually executing the trades — eating up all the apparent return?

The programming phase translates these rules into code that can be executed by a computer. In professional systematic trading operations, the code is typically reviewed by multiple people, tested in isolated environments, and subjected to extensive checks before it touches real money.

Testing: Looking Before You Leap

Testing a systematic strategy means running it on historical data to see how it would have performed. This is called a backtest. A backtest is a simulation of what the strategy would have done if it had been running in the past.

Backtesting sounds straightforward, but it is one of the most technically and intellectually demanding parts of systematic strategy development. The key is that the backtest must be honest. Every assumption in the backtest must reflect reality: transaction costs must be realistic, the timing of data availability must be accurate, and the historical prices must reflect what a real investor could have actually bought or sold at, not just the last printed price.

A good backtest shows not just the returns, but the full statistical profile of the strategy: how volatile the returns were, how severe the worst losses were (the drawdowns), how long those bad periods lasted, and how the strategy performed in different market environments.

A backtest is a necessary condition for taking a strategy seriously. It is not a sufficient condition. The next chapter discusses the additional steps required to ensure that what worked in the past has a genuine chance of working in the future.

CHAPTER FIVE

Building a Systematic Strategy — Part Two: Making It Real

In the previous chapter, we followed a systematic strategy from its initial idea through historical testing. Now we address the harder questions: how do we know if what we found is real? How do we make sure the strategy will behave in the future as well as it did on paper? And how do we understand whether it belongs in a real portfolio?

Optimisations: The Art of Disciplined Refinement

After the initial backtest, a developer will typically want to refine the strategy. Perhaps the core idea works, but the specific parameters — how many assets to hold, how often to trade, how aggressively to size positions — can be improved. This process is called optimisation.

Optimisation is useful and also dangerous. Useful, because finding better parameters is a legitimate way to improve a strategy's risk and return characteristics. Dangerous, because the more you

adjust parameters to fit historical data, the more you risk producing a strategy that is perfectly tuned to the past and useless for the future.

The discipline required here is to optimise as little as possible, to use as few free parameters as possible, and to accept that the best historical result is almost certainly too good to be true. Professional systematic developers often deliberately accept slightly worse historical results in exchange for a more robust strategy that uses simpler rules and fewer parameters.

Think of it this way: a recipe with ten ingredients, all carefully measured to seven decimal places, might taste perfect in one specific kitchen on one specific day. A recipe with five ingredients and simple proportions can be cooked reliably anywhere. The simpler recipe is the one you actually want.

Walk-Forward Testing and Randomisation

One of the most important tools for testing whether a systematic strategy is genuinely robust is the walk-forward test. Here is how it works.

Instead of using all the available historical data to build and test the strategy, a walk-forward test divides the data into two periods. The strategy is built using the first period (the in-sample period) and then tested on the second period (the out-of-sample period). The out-of-sample period was never used in building the strategy, so it provides a much more honest assessment of how the strategy will behave in real life.

This can be extended into a rolling walk-forward test: the strategy is repeatedly built on one window of data and tested on the next window, rolling forward through time. The resulting out-of-sample performance record gives a realistic picture of what live trading might look like.

Randomisation tests take this further. A developer can test whether the strategy's performance is genuinely better than random chance by creating many random variations of the strategy and comparing their results to the real strategy. If the genuine strategy

consistently outperforms thousands of random alternatives, that is meaningful evidence. If it barely beats the random alternatives, the developer should be very suspicious.

Statistical Analysis: Alpha, Beta, and What They Mean

To understand whether a systematic strategy is genuinely adding value, we need two concepts from financial statistics. Do not be alarmed — we will keep this simple.

Beta is the portion of a strategy's returns that can be explained by its exposure to general market movements. If a strategy simply holds a portfolio of stocks, its returns will largely move up and down with the stock market. That movement is its beta. Beta is widely available — you can get it cheaply by buying an index ETF.

Alpha is the return that cannot be explained by market exposure. It is the return that comes from the strategy itself — from its rules, its signals, its discipline. Alpha is what you are actually paying for when you hire an active manager or invest in a systematic fund.

A careful statistical analysis of a systematic strategy will separate these two components. A strategy that simply buys stocks when markets are rising and sells them when markets are falling has high beta but potentially no alpha. A strategy that generates positive returns regardless of whether markets are rising or falling — a truly market-neutral strategy — may have significant alpha with little beta.

The Sharpe ratio is the most commonly used measure to assess a strategy's quality. In plain terms, it measures how much return you get for each unit of risk you accept. A Sharpe ratio above 1.0 is generally considered good. The legendary Medallion Fund reportedly operated with Sharpe ratios well above 2.0, which is why it attracted so much attention from academics who could barely believe the numbers.

The Strategy Profile: What to Look For

By the time a serious strategy developer is ready to consider using a strategy with real capital, they should have assembled a comprehensive picture of what the strategy looks like across different market environments. This profile includes:

- The average annual return, stated realistically after all estimated costs.
- The volatility of returns — how much the returns vary from month to month and year to year.
- The maximum drawdown — the largest peak-to-trough decline in the strategy's value. This is the number that tells you how much pain you would have had to absorb during the worst period.
- The time to recovery — how long it took to recover from the worst drawdown.
- The Sharpe ratio — return per unit of risk.
- Correlation to equity markets and bond markets — how independently does the strategy behave relative to standard asset classes?
- Performance in specific historical stress periods — 2008, 2000-2002, 1997 Asia crisis, 1994 bond crash. How did the strategy behave when things went wrong?

No strategy will look perfect across all these dimensions. A strategy with very high average returns will almost certainly have high volatility and large drawdowns. A strategy with very small drawdowns will likely have modest returns. The art is in finding a strategy with a risk/return profile that complements your existing portfolio and matches your risk tolerance.

Portfolio Requirements for Adding a Strategy

A systematic strategy does not exist in isolation. When you are considering adding one to your portfolio, the most important question is not how good is this strategy on its own? The most important question is: how will this strategy behave together with what I already own?

A strategy that has modest returns but moves completely independently of your existing portfolio — low correlation — may add more value than a strategy with higher returns but that crashes exactly when your other investments crash. Diversification is not about owning many things; it is about owning things that do not all go wrong at the same time. We will come back to this principle repeatedly.

Practically, adding a systematic strategy to a portfolio requires consideration of minimum investment sizes, liquidity terms (how quickly can you access your money), fee structures, tax treatment, and the legal and regulatory framework of the investment vehicle. These are not trivial considerations, and for high-net-worth investors, the practical mechanics of access deserve as much attention as the underlying strategy.

CHAPTER SIX

Artificial Intelligence in Systematic Strategies: The Truth

The phrase artificial intelligence currently functions as a kind of magical incantation in the investment industry. Say it in a meeting and eyes widen, due diligence thickens, and allocations are discussed with new seriousness. As with most magical incantations, it is worth understanding what is actually behind the spell.

Data Mining: How Patterns Are Found

Artificial intelligence in the context of investment strategies is, at its core, about data mining — the systematic search for patterns in large datasets. If you have access to enough historical financial data and enough computing power, you can search for statistical relationships between different variables and the future returns of assets.

This is a legitimate and powerful capability. The question is not whether patterns exist in financial data — they do. The question is whether the patterns that are found by AI tools are genuine, persistent phenomena or statistical accidents that happened to appear in a particular dataset.

The history of quantitative finance is littered with strategies that looked extraordinary on historical data and failed spectacularly in live trading. The failure was not a problem with the data or the computing power. It was a problem of overfitting: finding patterns so specific to the historical data that they had no predictive power for the future. AI tools, precisely because they are so powerful at finding patterns, are also particularly susceptible to this problem.

Why Machine Learning Models Struggle with Prices

Traditional machine learning models have a fundamental problem when applied directly to financial price prediction: prices are extremely noisy, and the signal — the genuine, persistent, predictable component — is tiny relative to the noise.

In other domains where machine learning has been spectacularly successful — image recognition, natural language processing, medical diagnosis — the signal is relatively strong. A photograph of a cat contains an enormous amount of information about being a cat. A photograph of tomorrow's stock price contains almost no reliable information about what the stock price will be tomorrow.

The academic evidence on machine learning models applied directly to price prediction is sobering. In controlled studies where researchers carefully avoid data mining problems, the improvement in predictive accuracy from sophisticated AI models over simple traditional models is typically modest. Some studies find no statistically significant improvement at all.

This does not mean AI is useless in finance — far from it. It means that the application of AI requires careful thinking about where it can genuinely add value.

How Machine Learning Is Actually Used Well

The most sophisticated systematic investment operations use machine learning not as a magic price-prediction machine, but as a sophisticated statistical tool for examining the relationship between specific, economically motivated variables and future returns.

Here is the approach in plain terms. Instead of feeding raw prices into a machine learning model and asking it to predict tomorrow's price, a quantitative researcher will first create a set of indicators — variables derived from prices and other data — that they believe might carry information about future returns. These indicators might include measures of price momentum, changes in trading volume, corporate earnings trends, interest rate movements, or many other factors.

These indicators are sometimes called predictors or features in the machine learning literature. The machine learning model is then used not to predict prices directly, but to examine which predictors carry the most information about future returns, and how strong that relationship is.

This approach combines the judgment of experienced researchers — who choose which predictors to construct, based on economic reasoning — with the statistical power of machine learning, which can find complex relationships among many variables simultaneously. It is a partnership between human intelligence and computational power, not a replacement of one by the other.

Using AI Signals in Practice

When a machine learning model identifies a relationship between a predictor and future returns, that relationship is translated into a signal: a quantitative indicator that can be used to tilt a portfolio toward or away from a particular asset.

Importantly, no single signal is trusted blindly. A well-designed systematic system will combine many signals, weight them according to their historical reliability and statistical significance, and subject the combined signal to risk management rules before it translates into actual trades. The diversity of signals — using many different predictors that carry independent information — is a critical safeguard against any individual signal failing.

The most honest summary of the state of AI in systematic investing is this: it is a valuable tool that has improved the sophistication of systematic strategies in meaningful ways, but it has not changed the fundamental economics of the game. The underlying requirement — finding genuine, persistent, explainable market inefficiencies — remains as demanding as ever. AI is a better magnifying glass. But if there is nothing there to see, it will show you nothing more clearly than ever before.


CHAPTER SEVEN

Risk Management: The Only Thing That Keeps You Alive

There is an old saying in trading: being right about the market is not enough. You also have to be right about how much to bet. The history of finance is full of people who were right about the direction of a market but still went broke because they bet too much and could not survive the noise on the way to being proven correct.

Risk management is not a feature of systematic investing. It is the foundation. A systematic strategy without rigorous risk management is not a strategy — it is a gamble with a complicated backstory.

Plate III



The Mathematics of Loss

◇
a fall of half demands a doubling to return

Yes, Risk Management Must Be Systematic

Let us address the chapter's core thesis directly: yes, proper risk management must be systematic. Here is why.

Discretionary risk management — relying on a human being's judgment to decide when to reduce risk and when to increase it — faces a fundamental problem: humans are worst at exercising judgment precisely when their judgment matters most. During a market crisis, when emotions are running high and losses are mounting, the tendency to make poor decisions is strongest. Fear causes people to sell at the worst possible time. Overconfidence causes people to hold positions far too long. The whole arsenal of cognitive biases that be-

havioural economists have documented — loss aversion, anchoring, the disposition effect — all become most dangerous under stress.

A systematic risk management framework removes the human from this equation. The rules decide how much risk to take. The rules decide when to reduce positions. The rules decide when to stop trading a strategy that appears to be failing. No panic, no greed, no rationalisation. Just the rules.

The Problem with Pure Mathematical Approaches

Here is the important caveat: systematic does not mean that mathematical formulas can be blindly trusted without judgment.

Portfolio optimisation models — mathematical tools that calculate the theoretically best allocation of capital across a set of investments — can produce wildly inappropriate results if used without experienced oversight. The mathematical models that tell you the optimal portfolio weight for each asset are based on estimates of future returns, future risks, and correlations between assets. These estimates are uncertain, and small errors in the estimates can translate into enormous errors in the recommended portfolio weights.

A well-known example: a portfolio optimisation model given a slight difference in expected returns might recommend putting 80% in bonds and 20% in equities, or alternatively 20% in bonds and 80% in equities. The mathematical output is highly sensitive to inputs, but the inputs are highly uncertain. Trusting the model's extreme outputs without applying experienced judgment is a recipe for disaster.

The role of experienced risk managers and traders is precisely to interpret the mathematical outputs through the lens of practical reality, and to override recommendations that, while mathematically coherent, are obviously unreasonable. A model that recommends 100% concentration in one asset may be producing valid mathematics but nonsensical investment advice.

Concentration: The Silent Killer of Portfolios

One of the most common risk management failures is concentration — having too much exposure to a single asset, sector, geography, or risk factor. Concentration is dangerous because it means that if you are wrong about that specific thing, the consequences are severe.

Here is the critical distinction: owning 100 stocks does not guarantee diversification. If 80 of those 100 stocks are technology companies, you have a technology portfolio that happens to have many names in it. When the technology sector corrects — as it will, as it always has — all 80 positions will likely fall together. You thought you were diversified. You were not.

True diversification comes from diversifying risk, not diversifying names. The relevant question is not how many positions do I own? but rather how many genuinely independent sources of risk do I have? A well-designed systematic strategy pays careful attention to the correlations between positions — how closely their returns are linked — and manages concentration at the level of risk factors rather than simply at the level of individual positions.

Is Everyone Doing the Same Thing?

A reasonable concern about systematic investing, particularly the most popular strategies, is that they become crowded. If many systematic funds are using similar momentum-based rules to trade similar markets, might they all hit the exit at the same time, amplifying losses rather than dampening them?

This is a genuine risk, and it deserves an honest answer: yes, crowding happens, and yes, it can cause short-term disruptions. In 2007, a period referred to as the quant quake, several large quantitative equity funds simultaneously reduced their positions, causing a brief but violent dislocation in the prices of stocks commonly held by systematic strategies.

However, several factors mitigate this risk. First, systematic strategies are not all doing the same thing. The range of strategies — momentum, value, carry, mean reversion, macro — operates on

different signals, different timescales, and different markets. Their positions are often quite different. Second, the markets in which systematic strategies trade are enormous — global bond markets, currency markets, and commodity futures markets dwarf the capacity of even the largest systematic funds. Third, if a strategy truly becomes overcrowded, the returns available from it diminish, which naturally reduces the amount of capital attracted to it.

The crowding concern is real but manageable. The key for an investor is to ensure that the systematic strategies they invest in are not all using identical signals on identical markets. Diversity at the strategy level is the solution to the crowding problem.

CHAPTER EIGHT

The S&P 500 Fallacy: Why America Is Not the World

If you have spent any time in financial conversations over the past decade, you have encountered a particular argument with the tenacity of a persistent houseguest: the S&P 500 is the best investment in the world. Just buy it, hold it, stop overthinking.

There is genuine wisdom buried in this argument. There is also a significant amount of lazy thinking. Let us separate the two.

What the S&P 500 Actually Is


The S&P 500 is an index of 500 large companies listed on American stock exchanges, weighted by their market capitalisation (the total value of all their shares). When you buy an S&P 500 index fund, you are making the following bets simultaneously, whether you know it or not:

- That the United States will remain the world's dominant economy.
- That the stock market will remain the best-performing asset class within that economy.

- That the technology sector — which, despite the index's apparent breadth, represents a disproportionately large share of its value — will continue to grow.
- That large-cap companies will outperform smaller companies.
- That dollar-denominated assets are the right currency exposure for your wealth.

These are not unreasonable bets. Several of them have been excellent bets over the past decade. But they are bets, not certainties, and a serious investor should make them consciously rather than accidentally.

Plate IV



America Is Not the World

— ♦ —
one market among many, each with its own weather

The Historical Returns: Genuinely Good, But Context Matters

The historical returns of the S&P 500 are, by the standards of most asset classes, impressive. Over long periods, the index has generated average annual returns of roughly 7-10% in nominal terms, depending on the period chosen. For a passive investment requiring no research, no active management, and very low fees, this is an extraordinary result.

But context matters. The last century of American stock market history took place against a backdrop of American economic dominance that was historically exceptional. The United States emerged from World War II as the world's largest economy, with its industrial infrastructure intact while much of Europe and Asia was in ruins. It enjoyed decades of demographic growth, technological leadership, and relative political stability. These are not conditions that can be assumed to repeat indefinitely.

Furthermore, the returns of the S&P 500 are often quoted without adequate attention to the losses. The index fell approximately 50% from its peak during the dot-com crash of 2000-2002. It fell approximately 50% again during the global financial crisis of 2008-2009. After each crash, it eventually recovered and went on to new highs — but the road back took years. Someone who invested at the wrong point in time could have waited fifteen years to return to their original investment value in real terms.

The Concentration Problem

Here is the statistical reality that few discussions of the S&P 500 make explicit: the index is heavily concentrated in a small number of companies and a small number of sectors.

As of recent years, the five largest companies in the index have represented approximately 20-25% of its total value. The technology sector has represented roughly 25-30% of the entire index. This means that when you buy an S&P 500 diversified portfolio, you

are in reality placing a significant bet on a handful of American technology giants.

This is not diversification in any meaningful sense. This is a technology bet with equity characteristics, dressed up in the language of broad market exposure. The S&P 500 is a good investment vehicle. It is a poor definition of a diversified portfolio.

Real Diversification: Risk Across Countries, Assets, and Factors

Genuine diversification requires exposure to multiple independent sources of return. In practical terms, this means:

Geographic diversification: European equities, Asian equities, emerging market equities, and other developed markets all have their own economic cycles, political dynamics, and valuation characteristics. They will not all perform identically at the same time.

Asset class diversification: Bonds, real assets, commodities, and alternative strategies all have different return profiles that complement equities. A portfolio that holds only equities is not diversified — it is a single-factor bet.

Factor diversification: Value, momentum, quality, and other systematic risk factors carry different return streams that can be combined to create a more resilient portfolio.

The uncomfortable truth is that genuine diversification sometimes means owning things that are performing poorly while something else in your portfolio is doing well. An investor who abandons their international equity allocation because the S&P 500 outperformed for a decade has not diversified — they have anchored their portfolio to recent history and made it a reliable performance chaser.

Knowing the Risk You Are Taking

None of this is to say the S&P 500 is a bad investment. For many investors, particularly those with long time horizons and limited complexity preferences, it remains an excellent core holding. The

problem is not the investment — the problem is taking it without fully understanding the risk profile.

Before committing to any significant position — including the S&P 500 — an intelligent investor should be able to answer the following questions honestly:

- What is the maximum historical drawdown of this investment? (For the S&P 500: approximately -50% in 2008 and again in 2000-2002.)
- How long did it take to recover from the worst loss? (For the S&P 500: roughly 5-7 years from the 2000 peak in real terms.)
- What scenario might cause the investment to fail permanently or for a very long time? (For the S&P 500: prolonged US economic stagnation, regulatory reversal of technology dominance, or a prolonged period of high inflation.)
- Is my time horizon long enough to survive a recovery period of 5-10 years if necessary?

If you can answer these questions and still feel comfortable with the investment, then proceed. If you are surprised by the answers, that is a signal to reconsider the allocation.

A systematic strategy, when properly designed and evaluated, carries the same requirement for transparency about its risk profile. The next chapter will give you the tools to evaluate one.

CHAPTER NINE

Your Portfolio: Choosing and Sizing a Systematic Strategy

This chapter is the practical heart of the book. Everything we have discussed so far — the history, the mechanics, the risk framework, the diversification philosophy — comes together here. By the end of this chapter, you should have a clear framework for evaluating a systematic or alternative strategy and determining how it might fit into your portfolio.

Step One: Understand What Kind of Strategy You Are Looking At

Not all systematic strategies are the same, and the most important first step is to classify the strategy you are considering. This classification determines its role in your portfolio.

The most useful framework divides strategies into three categories, based on how their returns relate to the two main building blocks of most portfolios — stocks and bonds:

Equity-like strategies are those whose returns move broadly in the same direction as stock markets. Many investment funds fall into this category, even if they do not label themselves as such. A good diagnostic question: How did this strategy perform in 2008?

If the answer is badly, it likely has significant equity risk. This group includes most funds that trade individual company shares — whether buying and selling them to exploit price differences, betting on companies' bonds moving in value, trading currencies based on interest rate differences, or selling insurance-like contracts that pay out when markets are calm. All of these strategies tend to lose money when stock markets crash.

Bond-like strategies are those sensitive to interest rate movements. Strategies that invest in various types of lending — from corporate loans to infrastructure financing — fall here. They provide income and relative stability but tend to suffer when interest rates rise sharply.

Genuine alternatives are the holy grail: strategies with returns that are largely independent of both stock and bond markets. These provide genuine diversification — they add to your portfolio without simply duplicating existing risk.

The two most important categories of genuine alternatives are managed futures strategies and global macro strategies.

Managed futures strategies trade standardised contracts — called futures contracts — across a wide variety of markets: agricultural and energy commodities, currencies, government bonds, and stock market indices. Because they can bet on prices falling just as easily as rising, and because they trade across many unrelated markets simultaneously, they tend to generate returns that move differently from a traditional stock portfolio. During the worst stock market crashes in recent history, many managed futures strategies generated strongly positive returns. This is not a guarantee, but it reflects a genuine structural feature of how these strategies work.

Global macro strategies take broad views on major economic trends — the direction of interest rates in different countries, currency movements, commodity cycles — and express those views through a variety of instruments. They are more judgment-driven than pure managed futures strategies, but the best of them similarly tend to generate returns that do not simply track the stock market.

Step Two: Evaluate the Track Record Properly

Track records are the evidence you use to assess a strategy. They are also the most commonly abused piece of information in the investment industry. Here is how to evaluate one properly.

The most important principle: a track record only tells you something useful if it is long enough and the outperformance is large enough to distinguish genuine skill from luck. Financial markets are noisy. Short runs of good performance tell you almost nothing — a dice thrower who rolls six three times in a row has not discovered a new skill. They have had a lucky run.

How long is long enough? The uncomfortable answer is: longer than most people think. For a strategy with typical return characteristics, you generally need at least five to seven years of live track record before the evidence begins to be statistically meaningful. For strategies with higher variability or lower correlation to benchmarks, you may need longer still.

And the outperformance required is larger than most people assume. Even if a fund has genuinely better skill, it needs to demonstrate outperformance that is large enough relative to the variability of that outperformance to be convincingly distinguishable from chance. Academic research shows that for strategies with typical correlation to their benchmark, you would need to see consistent outperformance of several percent per year over many years before you can be genuinely confident the difference reflects skill rather than luck. A fund with three years of great returns is telling you almost nothing statistically meaningful.

Step Three: Look Past the Numbers to the Process

Here is a nuance that even sophisticated investors sometimes miss: for systematic strategies particularly, the track record is less important than the process that generates it.

A systematic strategy built on robust economic reasoning, thoroughly tested methodology, disciplined risk management, and experienced operators will continue to perform based on the same logic

that produced its historical results — as long as the fundamental market conditions that support the strategy remain intact.

A systematic strategy built on overfitted historical patterns — parameters tuned too precisely to past data, signals that have no economic justification, risk management that was never stress-tested — may have an excellent historical track record and a thoroughly unreliable future.

The process questions to ask when evaluating a systematic strategy include:

- What is the economic logic behind the strategy? Why should this signal predict future returns?
- How long is the out-of-sample track record, and does it match the in-sample performance?
- How many parameters does the strategy use, and were they optimised on the same data used to test the strategy?
- How did the strategy perform in 2008, 2020, and other stress periods?
- How does the team manage risk when the strategy is losing money? What are the rules for stopping trading or reducing risk?
- What is the maximum drawdown the strategy has experienced, and how was it managed?

Step Four: Size the Allocation Correctly

Assuming you have found a systematic or alternative strategy that passes your evaluation, the next question is: how much of your portfolio should it represent?

The key insight is that allocations should be sized based on risk, not simply on the amount of money involved. Because different strategies have different levels of variability in their returns — some swing widely, some move gently — putting the same dollar amount

into two different strategies does not mean you are taking the same risk. A highly volatile strategy with 10% of your capital might contribute as much risk to your portfolio as a calm, stable strategy with 30% of your capital.

Think of it this way: if you are building a portfolio and you want each strategy to contribute roughly equally to the overall risk you are taking, you will naturally put less money into the wilder strategies and more into the calmer ones. This is risk-based sizing, and it is the right way to think about it.

For most portfolios, a sensible allocation to genuine alternatives — strategies with genuinely independent return streams — is in the range of 10-15% of your total risk budget. In practical terms, depending on the strategy's volatility, this might translate to anywhere from 5% to 20% of the total portfolio value.

For a high-net-worth investor, a practical approach is:


- Start with a small allocation: 5-10% of total portfolio. This is enough to have a meaningful impact on portfolio diversification without creating excessive dependence on any single strategy.
- Scale up only after observing the strategy's live behaviour for at least twelve to eighteen months in your own portfolio.
- Maintain multiple strategies if possible: two to four genuine alternatives with different styles is better than a larger allocation to one.
- Do not allocate more than 20% of your portfolio to any single alternative strategy, regardless of how compelling its track record appears.

Step Five: Think About the Fit With Your Existing Portfolio

The most important question is not how good the strategy looks in isolation, but how it will behave in the context of your existing

holdings. This requires thinking about correlation — how closely does this strategy move together with what you already own?

A strategy with low or negative correlation to your existing portfolio — one that tends to go up when the rest of your portfolio goes down — is extremely valuable, even if its standalone returns are only moderate. This is what genuine diversification actually means, as opposed to the false diversification of owning many things that all go wrong at the same time.



Many Engines, One Voyage

◇
return streams that rise and fall out of step

As a practical check, look at the proposed strategy's monthly returns during the worst periods for your existing portfolio. The global financial crisis of 2008 is the most demanding test case. A

genuine alternative strategy should show significantly better returns during that period than your equity holdings. If it does not, you are almost certainly adding a strategy that will amplify your losses during the next crisis rather than offset them.

Step Six: Understand the Fees and Net Returns

Systematic and alternative strategies typically charge higher fees than passive index funds. This is sometimes justified — the best strategies do deliver performance that more than compensates for their fees — and sometimes not.

The standard approach for any active strategy is to evaluate net returns — returns after all fees, including management fees, performance fees, and any entry or exit costs. It is net returns that determine the actual benefit to your portfolio.

Be particularly careful about performance fees. A typical structure of 2% management fee and 20% performance fee means that a strategy generating 15% gross returns will deliver approximately 10% net after fees. If the strategy generates negative returns, you will still pay the 2% management fee. Performance fees have asymmetric consequences: in good years they significantly reduce your returns, and in bad years they do not refund your losses. High-water mark provisions — which require a fund to recover previous losses before charging new performance fees — are a minimum protection that you should require in any fund structure.

Step Seven: Consider Liquidity and Access

Systematic strategies, particularly those operating in less liquid markets or with complex approaches, often come with restrictions on when you can withdraw your capital. Redemption notice periods of 30, 60, or 90 days are common. Some strategies lock up capital for a full year or longer.

This is not inherently problematic — some of the most attractive returns in alternative investing come from accepting liquidity constraints. But you must ensure that the illiquid portion of your

portfolio does not exceed what you can comfortably lock away without affecting your financial flexibility. A general principle: do not put more capital in illiquid strategies than you could afford to leave untouched for the maximum expected lock-up period, even if unexpected circumstances required you to access capital.

For high-net-worth investors, a well-structured alternatives allocation might look something like this: a core of liquid, systematic strategies with monthly liquidity, combined with a smaller allocation to less liquid but potentially higher-returning opportunities where the illiquidity premium justifies the constraint.

CHAPTER TEN

Epilogue: Why Systematic Matters — and Why the Story Is Bigger Than That

We have come a long way from Chapter One, where we were merely establishing what the words mean. We have traced the history of systematic investing from its intellectual origins to its current sophistication. We have examined how strategies are built, tested, and deployed. We have looked at risk management with honest eyes. We have interrogated the S&P 500 myth. And we have given you a practical framework for evaluating and sizing systematic strategies in your own portfolio.

Let us close with a broader view.

Why Systematic Is the Future — but Not the Whole Story

The case for systematic investing is strong. Markets are increasingly competitive. Information spreads faster than ever. The advantages that discretionary managers once enjoyed — access to better information, faster analysis, superior data — have been progressively eroded by technology. In this environment, the disciplined, rule-based approach of systematic investing becomes ever more valuable.

Furthermore, as the volume of data available to investors grows — alternative data, satellite imagery, web traffic, credit card transaction data — the systematic approach is uniquely positioned to extract value from it. Human analysts cannot process thousands of data points simultaneously and update their views in real time. Systematic models can.

The best systematic funds of today are operating at levels of sophistication that would have been impossible a decade ago. The combination of better data, better computing power, and deeper understanding of how markets actually function has made the field significantly more capable. And this trend will continue.

So yes: systematic investing is the future of a significant portion of professional investment management.

But We Do Not Need to Choose a Camp

Here is the part that perhaps surprises readers who expected a book about systematic investing to end with a rallying cry for the systematic tribe.

We do not need to choose a camp.

Systematic investing and discretionary investing are not enemies. They are different approaches to the same fundamental challenge: understanding the world well enough to make better investment decisions than the next person. Both approaches, at their best, represent extraordinary human achievement — the disciplined application of intelligence to one of the most complex adaptive systems ever created by human beings.

The discretionary investor who can read a balance sheet, assess a management team's character, identify an emerging industry trend, and hold a position with patience through temporary adversity is doing something genuinely valuable. That investor, done right, is adding to the sum of human understanding about how capital is best deployed in the economy.

The systematic investor who builds a model grounded in rigorous research, stress-tests it relentlessly, implements it with disci-

pline, and manages its risk with precision is also doing something genuinely valuable — and complements the discretionary approach in ways that benefit markets overall.

Progress Is the Real Point

The deeper truth is this: humanity keeps moving forward. Our understanding of financial markets — their regularities, their anomalies, their structural features, their connections to the broader economy — is genuinely improving over time. The research is better. The data is better. The tools are better.

Whether that improvement comes through more sophisticated systematic models, deeper discretionary analysis, or some combination of the two, it contributes to the overall project of making capital allocation more rational, more efficient, and more reflective of genuine economic value. That is a project worth participating in.

For you, as a high-net-worth investor who has taken the time to read through this manual: the takeaway is not to become a systematic trader yourself, nor to abandon your existing approach entirely. The takeaway is to understand what systematic investing is, to evaluate it honestly and rigorously, to incorporate it into your portfolio in a way that genuinely improves your risk-adjusted returns over the long run, and to do so with the same discipline that the best systematic strategies bring to their own work.

Invest with your eyes open. Ask the hard questions. Demand honest answers. Do not pay for confidence when what you need is competence. And remember that the investment industry, for all its sophistication, is still largely in the business of selling you things — the responsibility for distinguishing what is genuinely valuable from what merely sounds impressive rests, ultimately, with you.

A Final Thought

The title of this book is *The Manual for the Systematic Investor*. But perhaps the most important quality it describes is not systematic in the technical sense — it is systematic in the human sense.

Disciplined. Consistent. Resistant to panic and overconfidence alike. Curious, evidence-driven, and humble enough to acknowledge what is genuinely uncertain.

Those qualities — discipline, consistency, humility, intellectual honesty — are what separate good investors from bad ones, and they are qualities that no particular strategy has a monopoly on.

The market will test you. It always does. The question is whether your approach can survive the test.

We hope this book helps you make sure it can.

Key Ideas, Illustrated

The following diagrams illustrate four concepts that are central to understanding systematic investing. Each visual is designed to build intuition rather than provide precise mathematical detail — consider them mental models to carry with you as you read.

Alpha vs. Beta: Two Sources of Return

Every investment return can be broken into two parts. Beta is the return you get simply from being exposed to the market — the tide that lifts (or sinks) all boats. Alpha is the additional return that comes from skill, insight, or a specific edge. The diagram below illustrates this separation.




Figure 1: Beta tracks the market; alpha is the return above and beyond it.

Walk-Forward Testing: In-Sample vs. Out-of-Sample

A strategy built and tested on the same data will always look better than it truly is — the model has seen the answers before the exam. Walk-forward testing solves this by keeping a portion of the data completely untouched during strategy construction. Only after the strategy is finalised is it tested on that unseen period.




Figure 2: The strategy is built on the In-Sample period and tested on the Out-of-Sample period it has never seen.

Rolling Walk-Forward Windows

The rolling version extends this idea across time. The in-sample and out-of-sample windows move forward together, producing a long out-of-sample track record that gives a realistic picture of what live trading would look like. This is one of the most rigorous ways to avoid the illusion of a great strategy that only works on paper.

EPILOGUE: WHY SYSTEMATIC MATTERS — AND WHY THE STORY IS
BIGGER THAN THAT




Figure 3: The windows roll forward in time, producing a continuous out-of-sample performance record.

Key Terms, Simply Explained

Every technical term used in this book is defined below in plain language. Terms appear in alphabetical order.

Alpha

The return generated by an investment strategy that cannot be explained simply by exposure to the overall market. If the stock market rises 8% in a year and your fund rises 12%, the 8% is beta (market exposure) and the extra 4% is alpha — the genuine value added by the manager or the strategy. Generating consistent, real alpha is extremely difficult, which is why so few active managers manage to do it over long periods.

Beta

A measure of how much an investment moves in line with the overall market. A portfolio with a beta of 1.0 tends to rise and fall in step with the market. A portfolio with a beta of 0.5 tends to move half as much as the market — less upside, but also less downside. Beta is the return you get just for participating in a market; you can capture it cheaply by buying a simple index fund. You do not need to pay a manager's fee for beta.

Book value

The accounting value of a company — essentially what you would theoretically get if you sold everything the company owns

(buildings, equipment, inventory, cash) and paid off all its debts. Book value gives a rough indication of what a company is worth based purely on its assets, without considering how profitable it might be or how much investors are willing to pay for its future growth. Value investors often look for companies trading at a low price relative to their book value.

Correlation

A measure of how closely two investments move together. Correlation runs from -1 to +1. A correlation of +1 means two investments always move in the same direction — owning both gives you no diversification benefit. A correlation of -1 means they always move in exactly opposite directions — ideal insurance, but rare in practice. A correlation of 0 means the two investments move independently of each other. The practical goal in portfolio construction is to combine investments with low or negative correlations so that when one falls, the others are not necessarily falling too.

Data mining

The process of searching through large amounts of historical data to find patterns. In finance, data mining is used to identify relationships between indicators and future returns. The danger is finding patterns that are purely accidental — patterns that appear in the historical data but have no real meaning and will not persist in the future. This is called overfitting. Good systematic strategy development uses rigorous testing to distinguish genuine, economically meaningful patterns from data mining accidents.

ETF (Exchange-Traded Fund)

A fund that holds a collection of assets — typically designed to track a market index — and is traded on a stock exchange just like an individual share. For example, an S&P 500 ETF holds shares in all 500 companies in that index in proportion to their size. ETFs are generally low-cost, transparent (you can see exactly what they hold), and easy to buy and sell. They are the practical vehicle most indi-

vidual investors use to get exposure to different markets and asset classes without having to buy every individual asset themselves.

Futures contracts

A legally binding agreement to buy or sell a specific asset at a predetermined price on a specific future date. For example, a futures contract might commit you to buying 1,000 barrels of oil at \$80 per barrel in three months' time, regardless of what the actual price is then. Futures are used by producers and consumers of commodities to lock in prices (a farmer locks in the selling price of their wheat harvest before it is harvested), and by investors to gain exposure to markets — or to bet on their direction — without holding the underlying asset directly.

Growth stocks

Shares in companies that are expected to grow their revenues and profits significantly faster than average, and which investors are willing to pay a premium price for today in expectation of that future growth. Technology companies are the classic example. Growth stocks typically have high price-to-earnings ratios — meaning you are paying a lot for every dollar of current earnings because you are really paying for future earnings. They tend to do well in periods of low interest rates and strong economic confidence, but can suffer badly during downturns or when expectations change.

Hedgers in commodity futures markets

Businesses or individuals who use futures contracts to protect themselves against price movements in commodities they produce or consume as part of their normal operations. A wheat farmer who sells wheat futures before harvest is a hedger — they are locking in a selling price and protecting themselves against the risk that the price falls dramatically before they can sell their crop. An airline that buys oil futures is hedging against the risk of fuel prices rising sharply. Hedgers are not trying to profit from market movements; they are trying to reduce uncertainty in their business. Systematic investors

often act as the counterparty to these hedgers, providing liquidity and, in return, potentially earning a return for bearing the price risk the hedger wants to avoid.

Index (such as the S&P 500)

A standardised measure designed to track the performance of a specific group of assets. The S&P 500 index tracks the 500 largest publicly listed companies in the United States, weighted by their market value. It is not something you can invest in directly — it is a measurement. What you can invest in is an ETF or fund that is designed to track that index as closely as possible. Other well-known indices include the FTSE 100 (100 largest UK companies) and the MSCI World (approximately 1,500 large and medium companies across 23 developed countries).

Liquidity provision to buyers and sellers

The act of being willing to buy when someone wants to sell, or sell when someone wants to buy, at a price that is close to the current market price. Providing liquidity makes markets function smoothly — without it, buyers and sellers would struggle to find each other and prices would become erratic. Some systematic strategies earn a return by consistently being willing to take the other side of trades from buyers or sellers who need to transact quickly, accepting a small unfavourable price in exchange for their willingness to trade immediately.

Managed futures fund

An investment fund that uses systematic, rules-based strategies to trade futures contracts across a broad range of markets — commodities, currencies, government bonds, and stock market indices. Because futures can be used to bet on prices falling as easily as rising, managed futures funds can potentially profit whether markets go up or down, as long as there are sustained directional trends. They are also known as Commodity Trading Advisors (CTAs). Managed

futures funds have historically shown low correlation to traditional stock and bond portfolios, making them valuable diversifiers.

Market microstructure

The detailed mechanics of how markets actually operate — how buyers and sellers find each other, how prices are set moment to moment, how orders are processed, and how trading costs arise. Understanding market microstructure is important for systematic investors because it affects the practical cost of executing a strategy. A strategy that looks excellent on paper may be significantly less attractive once you account for the costs of actually getting trades done in real markets, especially in less liquid markets where large orders can move prices against you before you finish executing.

Non-linear relationship

A relationship between two variables where changes in one do not produce proportional changes in the other. In a linear relationship, every additional unit of X produces the same change in Y. In a non-linear relationship, the effect of X on Y changes depending on the current level of X — it might be weak at first, then accelerate dramatically, or vice versa. Financial markets are full of non-linear relationships: the impact of a large trade on the market price is far more than proportionally larger than a small trade; extreme events occur far more frequently than simple linear models predict. Advanced machine learning tools can sometimes identify these non-linear patterns where traditional statistical methods would miss them.

Quantitative equity funds

Investment funds that use systematic, mathematical methods to select and weight holdings in company shares. Rather than relying on an analyst's judgment about a company's prospects, quantitative equity funds apply statistical models to rank thousands of companies on measurable characteristics — such as price momentum, valuation ratios, profitability trends, or analyst forecast revisions —

and build portfolios based on those rankings. They can process far more information simultaneously than any human analyst, but their success depends on identifying genuine, persistent patterns rather than temporary statistical noise.

Sharpe Ratio

The most widely used measure of an investment strategy's risk-adjusted performance. It answers the question: how much return am I getting per unit of risk I am taking?

$$\text{Sharpe Ratio} = \frac{\text{Average Return} - \text{Risk-Free Rate}}{\text{Volatility of Returns}}$$

The Average Return is the strategy's mean annual return. The Risk-Free Rate is the return you could earn with no risk at all — typically the yield on short-term government bonds. We subtract it because we only want to measure the return you are earning above what you could get for free. Volatility of Returns measures how much the returns fluctuate — a strategy that earns 10% every single year with no variation has lower volatility than one that earns 30% one year and -10% the next, even if the average is similar.

Why does this ratio matter? Because higher returns are not necessarily better if they come with dramatically higher risk. A strategy earning 15% per year with wild swings might actually be less attractive than one earning 10% per year with smooth, consistent returns.

As a rough guide: a Sharpe Ratio below 0.5 is weak; around 1.0 is solid; above 1.5 is excellent; above 2.0 is exceptional and quite rare over long periods. The legendary Medallion Fund reportedly operated with Sharpe Ratios consistently above 2.0 for decades — a record that remains essentially unmatched in the professional investment world.

Systematic macro funds

Investment funds that take views on large-scale economic themes — such as the direction of interest rates, the relative strength of

currencies, or commodity price cycles — using systematic, rules-based models rather than purely human judgment. They differ from managed futures funds in that they tend to focus more explicitly on macroeconomic factors and may use a wider variety of instruments beyond futures contracts. They differ from discretionary macro funds in that their decisions are driven by quantitative models rather than a single portfolio manager's reading of the global economy.

About the Author

John Panayotaras

John Panayotaras is the founder of Insight LLC, a firm focused on systematic and quantitative investment strategy. He holds a background in mathematics and has worked as a financial analyst, where he developed a hands-on understanding of how quantitative methods are applied in practice.

For a number of years he taught mathematics to professionals in the banking and computer science industries — many preparing for formal qualifications or career advancement — as well as to university students. That experience of making rigorous concepts genuinely accessible to working professionals is the spirit that runs through this book.

His work is grounded in the belief that intelligent investing requires both intellectual honesty about what works and the discipline to follow a process when markets are doing their best to make you abandon it.

The Manual for the Systematic Investor is his first book.

© 2026 Insight LLC. All rights reserved.

"For those who ask the harder questions."

This manual provides a framework for understanding what a systematic, rules-based strategy is. Starting from the basics and the history, then evolving to all the nuances, it offers the serious investor a disciplined, evidence-based approach to putting capital to work.

INSIDE THIS BOOK

The Language of the Game

The Mutual Fund Industry: A Cautionary Tale

The Origins of Systematic Investing

Building a Systematic Strategy

Artificial Intelligence in Systematic Strategies

Risk Management

The S&P 500 Fallacy

Your Portfolio

John Panayotaras is the founder of Insight LLC, specializing in systematic and quantitative investment strategies. With a background in mathematics and years of teaching professionals in finance and technology, he brings clarity and discipline to complex financial ideas.

INSIGHT LLC

© 2026 All Rights Reserved